

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾

Specialties ▾

Competences ▾

About Us ▾

Blog

Let's talk



Categories

Cybersecurity

Virtualization

Blockchain

C++ tips

Driver development

System development for Windows

Development for Linux

Development for Mac

Mobile development

Web Development

Reverse engineering

QA and testing

White Papers

Top content



Linux Driver Tutorial: I Simple Linux Device D

This Linux device driver tutorial will provide you with all the necessary information about how to write a device driver. This article includes a practical Linux driver development example that will follow. We'll discuss the following:

- Kernel logging system
- How to work with character devices
- How to work with user-level memory

We'll discuss how to write a Linux kernel version 2.6.32-5.3.el6.x86_64 driver. You can

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾

Specialties ▾

Competences ▾

About Us ▾

Blog

Let's talk



This article includes
description of simple



How to Reverse Engineer
Software (Windows) in a Right
Way

This article considers
common reverse engineering



Windows Process Monitoring
and Management Tips

The following article will help
you to understand principles

[Specifying a name of the device](#)

[The file_operations structure](#)

[The printk function](#)

[Using memory allocated in user mode](#)

[Build system of a kernel module](#)

[Loading and using a module](#)

[References](#)

1. Overview

Linux has a monolithic kernel. For this reason, it requires performing a combined compilation of the kernel and the driver. The first option is to implement your driver as a kernel module. In this case, you need to recompile the kernel to add another driver. The second option is to implement your driver as a net module.

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾ Specialties ▾ Competences ▾ About Us ▾

Blog

Let's talk



We run the module code in the kernel context, so we are very attentive, as it entails extra responsibilities when implementing a user-level application, and the user application in most cases; but if a developer is implementing a kernel module, the consequences are at a higher level. Luckily for us, the Linux kernel has a number of checks for errors in module code. When the kernel encounters an error (for example, null pointer dereferencing), you'll see a message. Some malfunctions during Linux operation are called "kernel oops". A malfunctioning module will be unloaded, allowing the system to work as usual. In addition, you'll be able to see a message that precisely describes this error. But be aware that a kernel panic message is not recommended, as doing so may lead to a system panic.

The kernel and its modules essentially represent a single program. Keep in mind that a single program module uses a single namespace. In order to minimize it, you must watch what is exported. Global characters must be named uniquely. A workaround is to simply use the name of the module as a prefix and must be cut to the first few characters.

Send us a message

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾

Specialties ▾

Competences ▾

About Us ▾

Blog

Let's talk



```
static void my_exit(void)
{
    return;
}

module_init(my_init);
module_exit(my_exit);
```

The only two things this module does is load driver, we call the *my_init* function, and to function. The *module_init* and *module_ex*. driver loading and unloading. The *my_init* a identical signatures, which must be exactly a

```
int init(void);
void exit(void);
```

If the module requires a certain kernel versio the version, we need to link the **linux/modul** module built for another kernel version will l prohibiting its loading. There's a reason for s API are released quite often, and when you c signature has been changed, you cause dama *module_init* and *module_exit* macros are

Send us a message

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾

Specialties ▾

Competences ▾

About Us ▾

Blog

Let's talk



development process richer.

For a start, here's some useful information a find device files in the /dev folder. They facil and the kernel code. If the kernel must receiv device file to pass it to the module serving tl device file originates from the module servin into two groups: character files and block file whereas block files are buffered. As their nar to read and write data character by characte write only whole blocks of data. We'll leave t the scope of this article, and will get straight

Linux systems have a way of identifying devi which identify modules serving device files o **device numbers**, which identify a specific de major device number specifies. In the driver c as constants or they can be allocated dynam constant has already been used, the system is allocated dynamically, the function reserve being used by anything else.

Send us a message

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾ Specialties ▾ Competences ▾ About Us ▾

Blog

Let's talk



which the device and the *file_operations* zero to the major parameter, the function will return the value it returns) on its own. If the value is zero, it signifies success, while a negative number signifies an error specified in the 0–255 range.

We pass the device name as a string value or can also pass the name of a module if it registers this string to identify a device in the `/sys/devices` directory. The `read`, `write`, and `save` are processed by the *file_operations* structure. These functions and the pointers to the *module* structure identify the device within the *file_operations* structure. Here is the version structure:

```
struct file_operations {
    struct module *owner;
    loff_t (*llseek) (struct file *, loff_t, int, unsigned int);
    ssize_t (*read) (struct file *, char *, size_t, loff_t *);
    ssize_t (*write) (struct file *, const char *, size_t, loff_t *);
    int (*readdir) (struct file *, void *, struct dirent **, int);
    unsigned int (*poll) (struct file *, struct pollfd *, int);
    int (*ioctl) (struct inode *, struct file *, unsigned int, void *);
    int (*mmap) (struct file *, struct vm_area_struct *);
    int (*open) (struct inode *, struct file *, unsigned int, void *);
    int (*flush) (struct file *);
    int (*release) (struct inode *, struct file *, int);
    int (*fsync) (struct file *, int, int, int);
    int (*fasync) (int, struct file *, int);
}
```

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾ Specialties ▾ Competences ▾ About Us ▾

Blog

Let's talk



unimplemented function can simply be set to 0, which will make the kernel take care of the implementation of the function. In our case, we'll just implement the *read* function.

As we're going to ensure the operation of our driver on Linux, our *file_operations* structure must be implemented. Correspondingly, after it's created, we'll need to register it. Here's how this is done:

```
static struct file_operations simple_driver_fo
{
    .owner    = THIS_MODULE,
    .read     = device_file_read,
};
```

The declaration of the *THIS_MODULE* macro is located in the *linux/module.h* header file. We transform the macro into the *THIS_MODULE* variable in the required module. A bit later, we'll get to the implementation of the *device_file_read* function, but right now we have only the prototype, but right now we have only the *device_file_read*.

```
ssize_t device_file_read (struct file *,
```

The *file_operations* structure allows us to

Send us a message

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾

Specialties ▾

Competences ▾

About Us ▾

Blog

Let's talk



```

        , device_file_major_number );
    return 0;
}

```

The *device_file_major_number* is a global device number. When the lifetime of the driver revoke the registration of the device file.

6. The printk Function

We've already listed and mentioned almost a *printk* function. The declaration of this function is in the `linux/kernel.h` file, and its task is simple: to log kernel messages. It has paid attention to the *KERN_NOTICE* and *KERN_WARNING* levels, which are present in all listed format strings of *printk*. The *WARNING* and *WARNING* signify the priority level of a message, from insignificant *KERN_DEBUG* to the critical *KERN_ALERT* and *KERN_EMERG*, which signify system instability. This is the only difference between *printf* library function.

The *printk* function forms a string, which will be printed to the kernel log. The kernel daemon reads it and prints it to the console.

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾

Specialties ▾

Competences ▾

About Us ▾

Blog

Let's talk



register_chrdev and the unregister_chrdev ways.

To register a device, we use the following code:

```
void unregister_device(void)
{
    printk( KERN_NOTICE "Simple-driver: unregi
    if(device_file_major_number != 0)
    {
        unregister_chrdev(device_file_major_nu
    }
}
```

7. Using Memory Allocated

The function we're going to write will read client data. The signature of this function must be appropriate. The *file_operations* structure:

```
ssize_t (*read) (struct file *, char *, size_t
```

Let's have a look at the first parameter, the *file* structure allows us to get information in

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾ Specialties ▾ Competences ▾ About Us ▾

Blog

Let's talk



...the kernel address space may have
cannot simply dereference the pointer. When
have a set of specific macros and functions to
file. The most suitable function in our case is
for itself: it simply copies specific data from
allocated in the user space. In addition, it also
the buffer size is large enough. Thus, errors
relatively easily. Here's the code for the *copy*

```
long copy_to_user( void __user *to, const void
```

First of all, this function must receive three parameters: the buffer, a pointer to the data source, and a function for copying. As we've mentioned, an error return case of successful execution, the value will be `0`. The `__user` macro, whose task is to perform documentation for a useful application that allows us to analyze the kernel address space correctly; this is done using the static code analysis of static code. Make sure to always use `__user`.

Send us a message

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾

Specialties ▾

Competences ▾

About Us ▾

Blog

Let's talk



```
    printk( KERN_NOTICE "Simple-driver: Device
bytes count = %u"
           , (int)*position
           , (unsigned int)count );
    /* If position is behind the end of a file
    if( *position >= g_s_Hello_World_size )
        return 0;
    /* If a user tries to read more than we ha
ave */
    if( *position + count > g_s_Hello_World_si
count = g_s_Hello_World_size - *positi
    if( copy_to_user(user_buffer, g_s_Hello_Wc
)
        return -EFAULT;
    /* Move reading position */
    *position += count;
    return count;
}
```

| 8. Build System of a Kernel

After we've written the code for the driver, it as we expect. In the earlier kernel versions (required many more movements from a deve compilation needed to be prepared individual the GCC compiler. Only after that would a de that could be loaded to the kernel. Fortunately the process is much simpler now. Today, muc makef t starts the kernel bu s an

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾

Specialties ▾

Competences ▾

About Us ▾

Blog

Let's talk



```
obj-m := module_name.o
module_name-objs := source_1.o source_2.o ... so
```

The *make* command initializes the kernel build

To build the module:

```
make -C KERNEL_MODULE_BUILD_SYSTEM_FOLDER M=`p
```

To clean up the build folder:

```
make -C KERNEL_MODULES_BUILD_SYSTEM_FOLDER M=`
```

The module build system is commonly locate

Now it's time to prepare the module build sy

execute the following command from the fol

located:

```
#> make modules_prepare
```

Finally, we combine everything and

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾

Specialties ▾

Competences ▾

About Us ▾

Blog

Let's talk



endit

The *load* target loads the build module and the kernel.

In our tutorial, we've used code from `main.c` driver. The resulting driver is named `simple-r`

9. Loading and Using Modul

The following command executed from the `s` the built module:

```
#> make load
```

After executing this command, the name of `t` `/proc/modules` file, while the device that the `/proc/devices` file. The added records look lik

```
Character devices: 1 mem 4 tty 4 ttyS ... 250 Si
```

The file `three` records contain `...` of t

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾

Specialties ▾

Competences ▾

About Us ▾

Blog

Let's talk



Hello world from kernel mode!

10. References

- 1. Linux Device Drivers, 3rd Edition by and [Greg Kroah-Hartman](#): <http://lwn.net>
- 2. The Linux Kernel Module Programn Ori Pomeranz: <http://tldp.org/LDP/lkmpg>
- Linux Cross Reference <http://lxr.free->

Download source code of [Simple Linux Driv](#)

We hope this tutorial comes in handy. You ca [development](#).

You may also be ir



Send us a message

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾

Specialties ▾

Competences ▾

About Us ▾

Blog

Let's talk



Get more posts l

Let's talk



Send us a message

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾

Specialties ▾

Competences ▾

About Us ▾

Blog

Let's talk



4000 chars left

Attach a file

Browse

I'm not a robot

reCAPTCHA
Privacy - Terms

By clicking Send you give consent
to processing your data

Send

[Portfolio](#)



[Case Studies](#)



[Blog](#)



[SDKs](#)



DMCA PROTECTED

Send us a message

Apriorit and selected trusted third parties use cookies on this site to improve performance, for analytics, and for re-marketing and re-targeting with relevant content. By browsing this site you are agreeing to this. For more information see our [Privacy Policy](#).

CLOSE



R&D Services ▾

Specialties ▾

Competences ▾

About Us ▾

Blog

Let's talk



Send us a message